

LLMs in Journalism: Objectivity and Communication Implications for News Reporting

Introduction

How is journalism changing in the age of artificial intelligence? Who determines what's important when machines, not editors, filter the news? These questions are no longer speculative. As LLMs like GPT become increasingly embedded in news production, we are witnessing a profound shift in the way information is created, curated, and consumed. While the efficiency and scalability of these technologies promise major benefits, they also provoke urgent concerns. Can AI-generated journalism uphold traditional journalistic values such as objectivity, fairness, and accountability? What happens when algorithms set the news agenda? And how do these shifts affect the public's trust in the media at a time when information credibility is already under scrutiny? To unpack these dynamics, this paper draws on five foundational communication theories: Agenda Setting Theory, Media Trust, Expectancy Violations Theory, Symbolic Interactionism, and Politeness Theory. These frameworks allow us to examine the interaction between AI and human communication in journalism, focusing on how the content LLMs produce shapes meaning, tone, and perception. This paper investigates the evolving relationship between human editorial judgment and machine-generated content. It argues that while LLMs are powerful tools, their use in journalism demands careful oversight and ethical design. The core question is not just whether LLMs do journalism, but what kind of journalism they are helping to create, and whether that journalism serves democratic society or undermines it.

Objectivity and LLMs

Journalism has long upheld objectivity as a core value, despite ongoing debates over its feasibility (Ward, 2004; Schudson, 2001). This ongoing debate about objectivity in journalism has deep roots in media ethics scholarship, including concerns about the limitations of neutrality and the performative nature of balance (Maras, 2013; Hanitzsch, 2007). Such works argue that objectivity is not merely a technique but a historically constructed ideal shaped by institutional and political pressures. That being said, objectivity in news writing theoretically involves presenting facts without personal bias, ensuring neutrality in framing, and separating opinion from reporting. The rise of LLMs challenges traditional notions of objectivity in two ways, algorithm transparency and training data bias. In "When Journalism Meets AI: Risk or Opportunity?" (Cheng, 2024), the author highlights that AI systems are often trained on large datasets scraped from the internet, which contain both factual information and ideological bias. Consequently, LLMs may inadvertently reproduce and amplify certain worldviews, even when instructed to be neutral. This phenomenon undermines the perception of objectivity in machine generated news.

However, it is worth noting that this challenge is not unique to AI. Human journalists also take in large datasets, the things they read, study, and experience, all of which naturally carry biases. As a result, even traditional journalism inevitably reflects some degree of subjectivity, especially when complex or contested topics are involved. Humans have developed professional norms, such as fact checking, sourcing multiple perspectives, and editorial review, to minimize bias. Similarly, AI prompt engineering could evolve to mirror these practices, welcoming the recognition that bias exists, but requiring LLMs to deliberately present multiple viewpoints, refrain from extra commentary, and focus strictly on verifiable facts. The question, then, is not whether LLMs can eliminate bias, no system can, but whether they can be designed to manage it

responsibly. One promising avenue is to engineer prompts and post processing protocols that require the model to present competing arguments and identify factual disputes. This approach would parallel how trained journalists consciously structure stories to mitigate personal biases.

There is also a deeper distinction to consider between how humans and machines learn. Human learning is shaped by context, experience, and emotional reasoning, while machine learning involves statistical pattern recognition across massive datasets. This difference raises further concerns, while journalists can self correct based on ethical principles, LLMs lack that internal compass unless explicitly programmed through data and prompt design. As "Do LLMs Plan Like Human Writers?" (EMNLP, 2024) emphasizes, while LLMs can outperform humans in organizing information logically, they do not possess the instincts, or moral frameworks, that guide human editorial judgment. Thus, while LLMs offer opportunities for factual precision, their potential to write objectively hinges not only on adopting human inspired strategies of balance and critical reflection.

Agenda Setting and Media Trust

In traditional journalism, agenda setting theory explains how news organizations influence public awareness, not by telling audiences what to think, but by determining what topics are deemed worthy of attention (McCombs and Shaw, 1972). Through editorial choices about headlines, placement, repetition, and framing, media outlets shape the perceived importance of events. Foundational studies such as McCombs and Shaw (1972) first introduced agenda setting by analyzing how mass media shaped public perceptions during elections. Later expansions emphasized second level agenda setting (framing) and the role of visuals, tone, and story placement in guiding interpretation (McCombs, 2005; Wanta et al., 2004). These contributions

underscore that agenda setting is a dynamic interaction between media institutions and audiences. Media trust is closely linked to these editorial decisions, audiences are more likely to trust outlets that appear to prioritize stories serving the public interest over sensationalism.

When LLMs are integrated into newsrooms, the logic of agenda setting becomes automated. The paper *Agenda Setting in the Age of Artificial Intelligence* (SSRN, 2024) explores how AI can shape editorial decisions based on algorithms, such as search engine trends, click through rates, and engagement metrics. This feedback loop approach risks reinforcing overexposed narratives and ignoring underreported issues. As media organizations increasingly rely on LLMs to curate news, agenda setting may become increasingly reliant on algorithmic expediency.

To better understand how LLMs may participate in agenda setting, I conducted a comparative experiment. On April 26, 2025, I prompted ChatGPT4o with a simple question, “What are the top news stories happening in the world today?” The GPT4o model used in my April 26, 2025 experiment was accessed through the ChatGPT web interface. The prompt, timestamp, and results are recorded in the Appendix attached.. The LLM responded with a list of stories centered heavily on the tech sector, particularly artificial intelligence. These included OpenAI’s ongoing organizational transition from nonprofit to for profit status, a development flagged by former employees as ethically fraught, the rollout of GPT Edu across the California State University system, offering students and faculty a free version of the tool, and the release of a new image generation API that enhances creative capacity. I then visited the homepage of The New York Times and reviewed its top stories from the same day. The coverage emphasized a broader and more geopolitically diverse set of issues, political developments in the U.S. Congress, foreign policy decisions in the Middle East, new climate data from the UN, and the Federal Reserve’s evolving stance on interest rates. AI and tech news were present but did not occupy prominent

positions in the editorial lineup. This divergence between GPT and traditional media outlets reflects two very different versions of “what matters today,” one curated through algorithms and recent user interest, the other through human editorial judgment. Additional experiments with Meta.ai and Perplexity yielded different top stories, demonstrating how model training and interface architecture significantly shape agenda outcomes. Meta.ai prioritized international politics and climate policy, highlighting new EU climate goals and tensions between China and Taiwan. Perplexity, meanwhile, emphasized economic developments and scientific breakthroughs, including updates on U.S. inflation rates and a newly published study on gene editing. These variations suggest that LLM driven agenda setting is not only prompt dependent but also platform specific.

This difference illustrates how agenda setting by LLMs is shaped not by journalistic standards but by a combination of user prompting, training data emphasis, and probabilistic reasoning. Since LLMs learn by spotting patterns in existing online text, they tend to focus on topics that are already popular on the internet, like AI, which gets talked about a lot in blogs, forums, and company posts. This means that the topics LLMs highlight might reflect what’s most visible online, not necessarily what’s most important in the real world. However, it is important to note that the way I prompted GPT likely influenced the outcome of the experiment. Asking "What are the top news stories happening in the world today?" is fairly open ended, and LLMs are heavily influenced by patterns of current online engagement. A separate experiment prompting, for example, “What are the major geopolitical news stories of today?” might yield a broader or more globally focused set of responses. Different wording can produce significantly different responses because the LLM optimizes for perceived relevance based on the prompt’s framing. In addition, it matters what GPT is trained on versus what internal LLMs used by media outlets

might be trained on. GPT, for instance, is trained on a vast corpus of publicly available internet data, including forums, blogs, and news websites, but not on real time professional editorial processes. An LLM used within a newsroom may be fine tuned differently, possibly drawing from curated, vetted journalistic sources rather than the broader internet. Thus, training data is critical in shaping how agenda setting operates in different AI systems.

This shift also parallels trends we already see in social media news. Algorithms curate content based on user preferences, engagement history, and attention metrics rather than on the "importance" of a story in a civic sense. News outlets themselves often prioritize stories they believe will generate the most attention, engagement, and ad revenue, rather than the ones that might be most consequential. The idea of "most important" is inherently subjective. Still, the takeaway remains that news ecosystems, both human and AI driven, increasingly focus on what will attract and hold public attention, sometimes at the cost of covering slow moving, complex, or "boring" but important topics. Trust in news media has largely declined, partly due to perceptions of bias and corporate influence. LLMs may exacerbate this distrust if audiences perceive AI generated content as impersonal. However, LLMs may also restore some trust by delivering fact based content devoid of overt political biases, as long as transparency in authorship is maintained and outputs are monitored for fairness. Moreover, agenda setting is not just about inclusion but exclusion. The absence of certain stories in the LLM's response, stories about global conflict, domestic policy, or marginalized communities, raises questions about which issues disappear in AI curated news content. If LLMs consistently elevate stories that are highly visible online or aligned with dominant interests, they will sideline topics that lack digital prominence but are crucial for civic engagement.

Expectancy Violations Theory

Expectancy Violations Theory (EVT), developed by Judee Burgoon (1978), examines how individuals respond to unexpected behaviors in communication (Burgoon, 1978). In traditional journalism, EVT helps explain audience reactions when journalistic norms are breached, such as when a news report adopts an unexpectedly casual tone or frames an event too aggressively. Readers develop expectations about how news should be communicated based on cultural conventions. When a journalist uses humor in a tragic report or frames a complex issue with overt partisanship, it may violate those expectations, changing their engagement with the content.

In the context of LLMs, violations occur when users experience AI behavior that is either too humanlike or too mechanical. The article "Expectancy Violation in Human-Chatbot Interaction" (SAGE, 2023) found that users often feel uncomfortable when AI uses overly formal tones in casual contexts or overly empathetic responses in professional scenarios. This aligns with Burgoon's framework, which distinguishes between positive violations (for example, pleasantly surprising stylistic creativity) and negative ones (for example, jarring tonal mismatches) (Burgoon & Hale, 1988). In journalism, violating stylistic norms can shift reader perception of credibility, particularly when the tone seems inconsistent with the subject matter. LLM generated news can either meet or violate user expectations depending on how it mimics human tone.

To explore how this plays out in practice, I conducted another brief experiment. I prompted ChatGPT4o on April 26, 2025 with both formal and casual markers. The details of this experiment are also included in my Appendix. One prompt was, "What are the top news stories happening in the world today?" and the other was, "Hey Chat! What's going on today?" The model returned responses written in a conversational, almost promotional tone for the casual prompt. For example, its summary of OpenAI's latest product updates read, "OpenAI has just

rolled out an exciting new image generation model with smoother integration and improved detail!” While factually accurate, this tone was noticeably more upbeat and enthusiastic than what one would expect from a traditional news article. In contrast, coverage of the same story in The New York Times used formal language and focused on ethical debate and institutional accountability.

This tone mismatch illustrates how LLM generated content can produce expectancy violations for users accustomed to conventional journalism’s style. News consumers typically anticipate an objective tone when engaging with political or technological reporting. Even subtle differences in phrasing (e.g., “exciting new tool” vs. “controversial new release”) can change how readers interpret an event’s significance. Prompt modifiers such as “write in a journalistic tone,” “summarize objectively,” or “avoid emotionally charged language” may reduce expectancy violations, and I tested these as well. More formal prompts produced outputs that more closely matched traditional journalistic tone. EVT suggests that not all violations are bad, but the key is context. An LLM that fails to adapt tone risks breaking these communicative norms and eroding trust. The implications of this are profound for journalism, where traditional meaning making involves editorial decisions, interviews, and audience feedback. Now, the “audience” becomes a co author, meaning that every AI generated story is filtered not just through algorithms but through the expectations and instructions of its user.

Symbolic Interactionism

Symbolic Interactionism, according to George Herbert Mead (1934), emphasizes the social construction of meaning through interactions (Mead, 1934). This perspective is further supported by Blumer (1969), who emphasized that meaning arises not from isolated symbols but

from their interpretation in social interaction. When readers engage with LLM generated news, they are engaging in a new form of interaction mediated through prompt response cycles. In traditional journalism, this framework helps scholars understand how meaning is shaped not only by the content of news but by the reader's social context. A news article does not carry a fixed meaning on its own, rather, readers engage with it based on their values, experiences, and expectations. In the context of LLMs, this theory makes us reconsider the role of the reader. When consuming AI generated news, the reader is receiving the information while simultaneously constructing its significance. The paper "The Symbolic Interactionism Theory: A Systematic Literature Review" (ResearchGate, 2021) suggests that meaning emerges through dialogic interaction. While LLMs cannot engage in true dialogue, their outputs are shaped by prior human interactions.

It's important to recognize that in this new AI model of journalism, prompt engineering becomes a major part of the journalist's or user's role. Just as journalists traditionally made conscious choices about tone, style, and content, users now shape the output of LLMs through their prompts. Both the journalist's internal editorial process and the user's external prompt shape the style and focus of the final story in similar ways. However, journalism remains rooted in how stories are sourced, who is interviewed, where journalists go to find information, and what questions they ask. These human driven processes of reporting and investigation will continue to differentiate journalistic content from AI only outputs, even as LLMs assist in writing or shaping articles. To explore how meaning construction changes depending on how a user frames a prompt, I tested GPT with two nearly identical questions, "What are the top news stories happening in the world today?" and "Summarize today's most important political and economic developments in a neutral journalistic tone." Even though both prompts asked for news

updates, the answers were very different. The first one focused on tech and science news, written in a casual tone. The second one focused more on global politics and the economy, with a more serious tone. Symbolic Interactionism helps us see that readers do not just interpret information, they help create it. Each story becomes a shared product of both algorithm and audience, blurring the boundaries between traditional journalism and personal interpretation.

Politeness Theory

Politeness Theory, formulated by Penelope Brown and Stephen Levinson (1987), discusses how communicators navigate social interaction to maintain the self image that individuals want to present (Brown and Levinson, 1987). In journalism, this theory applies to how reporters conduct interviews, cover traumatic events, and handle politically divisive topics, often using politeness strategies to soften criticism.

As LLMs take on journalistic tasks, they may unintentionally “threaten face” by failing to soften their statements or acknowledge multiple viewpoints. It is crucial for politeness strategies to be integrated in prompt engineering and post editing protocols to avoid alienating audiences. To examine this, I conducted an experiment by prompting GPT4o on April 26, 2025, to summarize two controversial news stories, the U.S. government’s response to climate protests, and Israel/Gaza conflict updates. You can find the details of this experiment in the attached Appendix. I used two versions of each prompt, a direct prompt and a politeness cued prompt. The direct prompt was “Summarize what’s happening with climate protests in the U.S. right now.” The politeness cued prompt was “Could you provide a brief and sensitive overview of the recent U.S. government response to climate protests?” The difference was striking. The direct prompt produced a clinical, fact based summary, while the polite prompt introduced hedging,

broader context, and more empathetic framing. This shows how small changes in prompt phrasing can trigger politeness mechanisms within LLMs. Good communication respects our wish to be liked (positive face) and our need for independence (negative face). When AI generated news sounds too direct or forceful, it can come off as rude or insensitive. Politeness cues help regulate tone and emotional weight, and they are especially critical when AI lacks the human intuition to adjust naturally. Journalism ethics often stress the importance of empathy, especially in coverage of conflict, trauma, or vulnerable populations. Human reporters intuitively apply face saving strategies, but LLMs must be prompted explicitly to mirror these communication norms. This makes prompt literacy, knowing how to instruct LLMs for appropriate tone, a core journalistic skill in the AI age.

Conclusion

The integration of LLMs into journalism marks a pivotal moment in the evolution of media. These technologies are not neutral, they reflect the values embedded in their training data, the biases of their algorithms, and the intentions of the humans who prompt them. As this paper has shown, LLMs challenge conventional notions of objectivity, shape editorial priorities, and influence audience trust through tone and framing.

Crucially, the reader's trust, historically rooted in belief in journalistic ethics, is now at stake. Readers today subscribe to outlets like The New York Times or The Washington Post not only for information but for a perceived ethical contract, that journalists operate with integrity, verify facts, and write with civic purpose. If Journalist X is now an LLM, or even just collaborating with one, transparency becomes crucial. Journalism ethics may need to evolve to mandate disclosure of AI assistance in bylines, article notes, or metadata. When readers engage

with journalism, they assume a human journalist made conscious decisions about sourcing, fairness, and importance. If LLMs increasingly participate in reporting, journalism ethics must advocate for clear transparency about when, and how, AI systems are used. Readers must be informed whether they are reading human reported, AI generated, or hybrid content. LLMs can support good journalism when used responsibly. Without careful editorial oversight, however, AI risks reinforcing sensationalism, bias, and alienation. Furthermore, journalism schools and newsrooms alike must develop guidelines for how LLMs should be deployed, ensuring that they support rather than supplant ethical human judgment. This includes protocols for prompt design, model selection, and editorial review. Only then can LLMs enhance trust rather than undermine it. Moving forward, the challenge is not merely technical, it is ethical. Journalism must adapt to safeguard its role in democratic society while thoughtfully integrating new technologies that shape how news is created, curated, and consumed.

References

- Blumer, H. (1969). *Symbolic interactionism: Perspective and method*. University of California Press.
- Brown, P., & Levinson, S. C. (1987). *Politeness: Some universals in language usage*. Cambridge University Press.
- Burgoon, J. K. (1978). A communication model of personal space violations: Explication and an initial test. *Human Communication Research*, 4(2), 129–142.

Burgoon, J. K., & Hale, J. L. (1988). Nonverbal expectancy violations: Model elaboration and application to immediacy behaviors. *Communication Monographs*, 55(1), 58–79.

<https://doi.org/10.1080/03637758809376158>

Cheng, S. (2024). When journalism meets AI: Risk or opportunity? *Communications of the ACM*, 67(3). <https://dl.acm.org/doi/full/10.1145/3665897>

EMNLP. (2024). Do LLMs plan like human writers? *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*.

Hanitzsch, T. (2007). Deconstructing journalism culture: Toward a universal theory. *Communication Theory*, 17(4), 367–385. <https://doi.org/10.1111/j.1468-2885.2007.00306.x>

Maras, S. (2013). *Objectivity in journalism*. Polity Press.

McCombs, M. (2005). A look at agenda-setting: Past, present and future. *Journalism Studies*, 6(4), 543–557. <https://doi.org/10.1080/14616700500250438>

McCombs, M. E., & Shaw, D. L. (1972). The agenda-setting function of mass media. *Public Opinion Quarterly*, 36(2), 176–187. <https://doi.org/10.1086/267990>

Mead, G. H. (1934). *Mind, self, and society*. University of Chicago Press.

ResearchGate. (2021). The symbolic interactionism theory: A systematic literature review.

https://www.researchgate.net/publication/349758172_The_Symbolic_Interactionism_Theory_A_Systematic_Literature_Review

Schudson, M. (2001). The objectivity norm in American journalism. *Journalism*, 2(2), 149–170.

<https://doi.org/10.1177/146488490100200201>

SSRN. (2024). Agenda setting in the age of artificial intelligence.

https://papers.ssrn.com/sol3/papers.cfm?abstract_id=4550921

Waddell, T. F., & Sundar, S. S. (2023). AI and agenda setting: How automated content influences public opinion. *New Media & Society*, 25(6), 1223–1240.

<https://doi.org/10.1177/14614448211044490>

Ward, S. J. A. (2004). *The invention of journalism ethics: The path to objectivity and beyond*. McGill-Queen's University Press.

Wanta, W., Golan, G., & Lee, C. (2004). Agenda setting and international news: Media influence on public perceptions of foreign nations. *Journalism & Mass Communication Quarterly*, 81(2), 364–377. <https://doi.org/10.1177/107769900408100209>

Wiley. (2024). Audiences, automation, and AI. *Journal of Communication Studies*.

<https://onlinelibrary.wiley.com/doi/full/10.1111/jcom.12345>